



Qwen3.5-9B Llama-3.1-8B Olmo-3.1-32B simulated


simulated: Llama-3.1-8B




simulated: Llama-3.1-8B




simulated: Llama-3.1-8B




simulated: Llama-3.1-8B




simulated: Llama-3.1-8B




simulated: Llama-3.1-8B



simulated: Llama-3.1-8B



simulated: Llama-3.1-8B



Trust-Elo by model and span — every number in the figures, and the run each came from

benevolence test split · newest judged sweep per model and span among the selected runs · 200-round bootstrap over rows, seed 0 · win rate is against each row's own unsteered reply, both slot orders pooled · 1500 is each sweep's own unsteered baseline, by anchor · SIMULATED series are placeholders, not measurements

Qwen3.5-9B · user · no penalty
vector run 20260913-125025_Qwen3.5-9B_benevolence_Lall_user
judged from data/trust-elo-2090156-20260913-125025.jsonl (4898 games)

condition	win %	95% CI	games	score	95% CI	graded	coherent	elo	95% CI
-2	30.9	[27.0, 35.2]	728	75.5	[70.5, 78.9]	400	93.0%	1360.5	[1327.9,1394.1]
-1	45.8	[41.8, 49.5]	746	85.1	[80.9, 88.8]	400	95.5%	1471.1	[1442.9,1496.3]
0	—		—	89.5	[86.0, 92.1]	400	96.8%	1500.0	[1500.0,1500.0]
+1	56.3	[52.2, 60.0]	748	91.3	[88.1, 93.8]	400	96.5%	1543.8	[1515.1,1570.3]
+2	66.4	[62.8, 69.8]	756	93.4	[90.8, 95.5]	400	97.0%	1618.2	[1590.6,1645.7]
distrust pole	25.5	[20.9, 30.2]	756	81.7	[77.1, 85.5]	200	97.5%	1304.1	[1254.1,1344.9]
trust pole	47.8	[42.7, 51.9]	774	95.6	[93.2, 97.8]	200	100.0%	1492.3	[1459.2,1521.3]

Llama-3.1-8B-Instruct · user · SIMULATED
NOT A MEASUREMENT: the mean of 6 measured sweeps, rescaled and jittered

condition	win %	95% CI	games	score	95% CI	graded	coherent	elo	95% CI
-2	36.9	[33.1, 41.1]	731	81.9	[77.5, 86.2]	388	93.4%	1394.6	[1361.6,1425.0]
-1	44.3	[40.7, 47.8]	752	85.9	[82.2, 89.3]	391	95.9%	1464.9	[1439.0,1489.7]
0	—		—	90.9	[87.5, 93.8]	391	96.5%	1500.0	[1500.0,1500.0]
+1	55.5	[51.9, 59.0]	755	90.6	[87.6, 93.4]	392	96.7%	1545.3	[1519.2,1570.3]
+2	62.3	[58.5, 65.8]	758	93.9	[91.3, 96.2]	394	97.3%	1604.4	[1574.2,1635.5]

0lmo-3.1-32B-Instruct · user · rpo 0.5
vector run 20260913-184027_0lmo-3.1-32B-Instruct_benevolence_Lall_user_rpo0.5
judged from data/trust-elo-2113270-20260913-184027.jsonl (4952 games)

condition	win %	95% CI	games	score	95% CI	graded	coherent	elo	95% CI
-2	34.0	[30.2, 38.3]	748	85.1	[80.6, 89.0]	382	95.5%	1384.6	[1357.6,1414.6]
-1	42.2	[38.5, 45.6]	756	85.5	[81.2, 89.2]	384	96.0%	1445.4	[1419.1,1470.2]
0	—		—	91.0	[86.9, 93.9]	388	97.0%	1500.0	[1500.0,1500.0]
+1	55.9	[51.8, 59.3]	768	92.4	[89.0, 94.8]	393	98.2%	1540.8	[1515.7,1564.0]
+2	68.1	[63.8, 71.5]	768	95.6	[93.4, 97.5]	394	98.5%	1631.5	[1600.3,1663.4]
distrust pole	35.6	[30.3, 40.8]	750	89.1	[85.5, 92.9]	193	96.5%	1389.5	[1342.9,1426.0]
trust pole	65.3	[61.1, 69.3]	776	95.5	[92.4, 98.2]	200	100.0%	1617.5	[1588.6,1646.2]

Trust-Elo by model and span — every number in the figures, and the run each came from (cont.)

Qwen3.5-9B · assistant · rpo 0.2

vector run 20260919-101411_Qwen3.5-9B_benevolence_Lall_assistant_rpo0.2

judged from data/trust-elo-60165-20260919-101411.jsonl (4870 games)

condition	win %	95% CI	games	score	95% CI	graded	coherent	elo	95% CI
-2	25.6	[21.9, 30.0]	698	76.2	[71.7, 80.4]	400	89.5%	1315.4	[1275.3,1346.3]
-1	41.6	[38.4, 44.9]	750	88.8	[85.8, 91.4]	400	95.5%	1441.2	[1415.8,1464.3]
0	—		—	90.6	[87.8, 93.4]	400	97.0%	1500.0	[1500.0,1500.0]
+1	56.2	[52.8, 59.9]	756	91.4	[88.6, 94.1]	400	96.8%	1543.4	[1521.0,1570.0]
+2	63.1	[58.8, 66.5]	750	93.9	[91.1, 96.3]	400	96.5%	1592.8	[1560.1,1624.8]
distrust pole	25.5	[21.7, 29.8]	752	83.2	[79.3, 86.7]	200	97.0%	1297.4	[1253.3,1335.3]
trust pole	46.9	[42.1, 51.0]	776	91.5	[88.0, 94.9]	200	100.0%	1490.8	[1457.8,1518.3]

Llama-3.1-8B-Instruct · assistant · SIMULATED

NOT A MEASUREMENT: the mean of 6 measured sweeps, rescaled and jittered

condition	win %	95% CI	games	score	95% CI	graded	coherent	elo	95% CI
-2	28.2	[24.4, 32.4]	731	80.6	[76.2, 84.9]	388	93.4%	1327.1	[1294.1,1357.5]
-1	40.0	[36.5, 43.5]	752	86.8	[83.1, 90.2]	391	95.9%	1428.9	[1403.0,1453.7]
0	—		—	89.2	[85.8, 92.1]	391	96.5%	1500.0	[1500.0,1500.0]
+1	58.2	[54.6, 61.7]	755	93.5	[90.4, 96.2]	392	96.7%	1558.8	[1532.7,1583.7]
+2	68.9	[65.1, 72.3]	758	95.3	[92.7, 97.5]	394	97.3%	1623.7	[1593.6,1654.9]

Olmo-3.1-32B-Instruct · assistant · rpo 0.3

vector run 20260922-080331_Olmo-3.1-32B-Instruct_benevolence_Lall_assistant_rpo0.3

judged from data/trust-elo-olmo32a-full-20260922-080331.jsonl (5069 games)

condition	win %	95% CI	games	score	95% CI	graded	coherent	elo	95% CI
-2	21.2	[17.8, 24.4]	764	77.6	[72.7, 83.1]	400	96.5%	1272.4	[1234.3,1303.4]
-1	37.3	[33.9, 41.1]	778	85.2	[81.4, 89.2]	400	97.8%	1409.7	[1384.1,1437.4]
0	—		—	89.6	[86.7, 92.6]	400	98.8%	1500.0	[1500.0,1500.0]
+1	62.6	[58.6, 65.9]	784	92.4	[89.5, 95.5]	400	99.0%	1589.6	[1560.5,1614.0]
+2	75.1	[71.3, 79.0]	788	93.8	[91.2, 96.6]	400	99.2%	1691.7	[1657.9,1729.9]
distrust pole	30.1	[25.1, 34.4]	777	77.2	[72.3, 81.7]	200	98.5%	1359.0	[1319.7,1394.2]
trust pole	69.1	[64.9, 73.6]	786	97.3	[95.1, 99.2]	200	99.5%	1634.6	[1605.9,1665.2]